

“Can we still use MD5?”

An md5 collision attack in practice


Simon Fransson

MD-what?

- “message-digest algorithm” v5
- Designed by MIT professor Ronald Rivest in 1992, when it was predicted MD4 would become insecure (spoiler alert, it did!).
- Computes a stable hash based on the contents of a file



my_file.txt



c263e623fb6bc74850a8dbfad7301af5

Can I still use MD5?

It depends.

The MD5 message-digest algorithm is a **cryptographically broken** but **still widely used** hash function producing a 128-bit hash value.

Different types of attacks

- **Collision attack** 🗝️ - find two different messages $m1$ and $m2$ such that $\text{hash}(m1) = \text{hash}(m2)$.
- **Chosen-prefix attack** 🗝️ - given two different prefixes $p1$ and $p2$, find two appendages $m1$ and $m2$ such that $\text{hash}(p1 \parallel m1) = \text{hash}(p2 \parallel m2)$,
- **Pre-image attack** 🗝️ - find a message m such that it produces an already known hash h , $\text{hash}(m) = h$

Super simple C program

```
#include <stdio.h>
```

```
int main(int argc, char *argv[])  
{  
    printf("Hello world!\n");  
    return 0;  
}
```

Looking at our executable

Looking at the compiled executable we can see the relevant byte code is prepended with a bunch of NULL bytes (macOS + x64, may vary by architecture):

```
gcc -o executable source.c
xxd executable | grep -B 50 Hello
00003ca0: 0000 0000 0000 0000 0000 0000 0000 0000  ....
00003cb0: 0000 0000 0000 0000 0000 0000 0000 0000  ....
00003cc0: 0000 0000 0000 0000 0000 0000 0000 0000  ....
00003cd0: 0000 0000 0000 0000 0000 0000 0000 0000  ....
00003ce0: 5548 89e5 4883 ec10 897d fc48 8975 f048  UH..H....}.H.u.H
...
00003f20: 00e9 aaff ffff 4865 6c6c 6f20 776f 726c  ....Hello worl
```

This is something we can exploit!


hashclash

[Project HashClash - MD5 & SHA-1 cryptanalysis](#)

“Project HashClash is a Framework for MD5 & SHA-1 Differential Path Construction and Chosen-Prefix Collisions for MD5. It's goal is to further understanding and study of the weaknesses of MD5 and SHA-1.”

Comes with a tool called `fastcoll` - fast MD5 collision generator.

Split executable into prefix and suffix



Split executable into prefix and suffix


Collision appendage is typically 128 (0x80) prepended with 32 (0x20) empty bytes, so let's find 160 unused bytes and use that to split our executable in two parts - **prefix** and **suffix**.

```
# addr      coll  pad    result
0x00003ce0 - 0x80 - 0x20 = 0x00003c40
```

```
dd if=executable bs=1 status=None count=$((0x00003c40)) of=prefix
dd if=executable bs=1 status=None skip=$((0x00003ce0)) of=suffix
```

Different prefixes, same MD5 hash

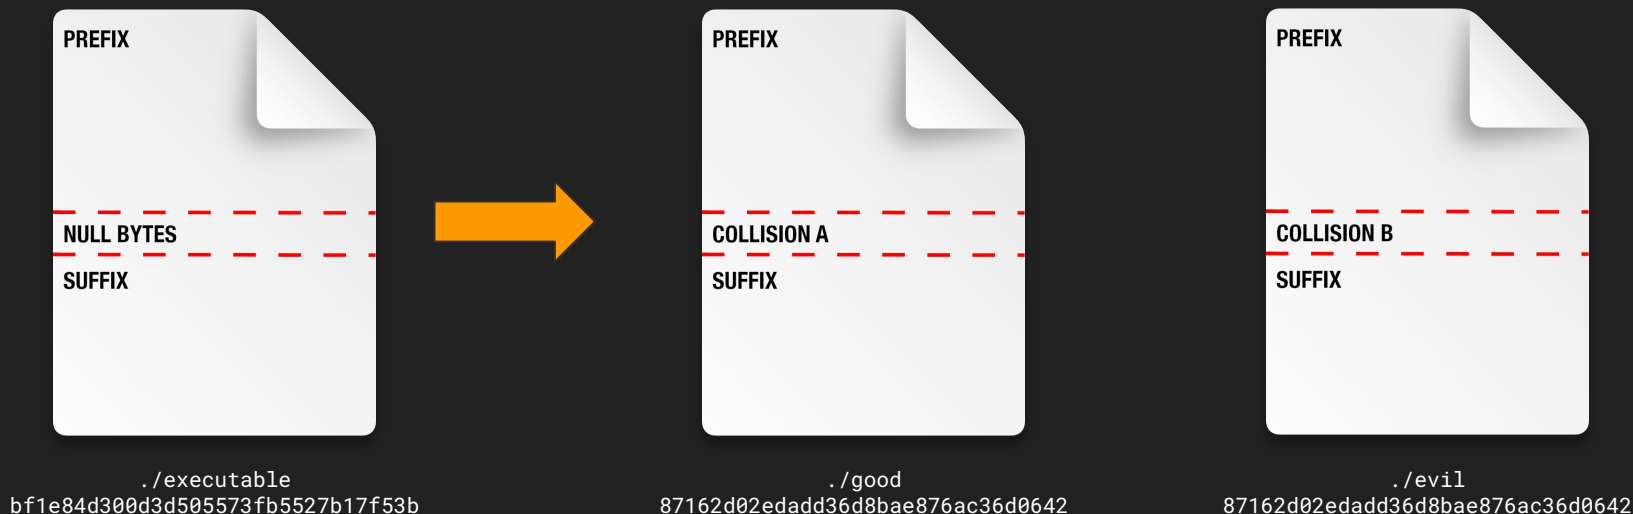
```
md5_fastcoll -p prefix -o prefix_a prefix_b
```



Join prefix and suffix

```
cat prefix_a + suffix > good
```

```
cat prefix_b + suffix > evil
```



Same program with a hidden payload

```
#include <stdio.h>

int good(int argc, char *argv[])
{
    printf("Hello world!\n");
    return 0;
}

int evil(int argc, char *argv[])
{
    // TODO: Be evil
}

int main(int argc, char *argv[])
{
    if (0) // TODO: Find a branch condition
    {
        return evil(argc, argv);
    }

    return good(argc, argv);
}
```

Condition for executing the evil payload

A diff between the two prefixes shows that they consistently differ in only a couple of bits!

We can use this to let the executable read itself and lookup the value at a specific position.

[illegible]

Payload branching added

```
#include <stdio.h>

int good(int argc, char *argv[])
{
    printf("Hello world!\n");
    return 0;
}

int evil(int argc, char *argv[])
{
    // TODO: Be evil
}

int main(int argc, char *argv[])
{
    if (read_byte(0x00000693) == 0xff) // TODO: How do we know which value to test against?
    {
        return evil(argc, argv);
    }

    return good(argc, argv);
}
```

Toggle evil payload

We can't just lookup the expected byte value and recompile + run fastcoll, as the value will change with each new build. But we can lookup the value and modify the compiled binary, or use a recompiled suffix part (which is easier than finding that one byte in a... byte stack?).

```
BYTE=$(xxd -s 0x00000693 -l 1 evil | awk '{ print $2 }')
sed -i '' 's/== 0xff/== 0x${BYTE}/' source.c
gcc -o executable source.c
dd if=executable bs=1 status=none skip=$((0x00000700)) of=suffix
cat prefix_col1 suffix > good
cat prefix_col2 suffix > evil
```

Demo time!

Implications

- Provide certain website visitors with a (seemingly) normal file and others with a malignant copy, both sharing the same MD5.
- Make adjustments to critical binaries (gcc etc.) on System A, while System B appears normal and they share the same MD5 ([Reflections on Trusting Trust](#)).
- Avoid MD5 (and SHA-1), SHA-2 is probably better

Thanks!

<https://github.com/dessibelle/md5-collision>